Annotation Interface

The web interface for each paper is at `http://localhost:port/arxiv_id`

Link to the Original Paper

Link to the Indexed Bibliography (in case you want to open two tabs side by side)

Notes

1. You have to click the 'submit' button to save data for **every individual table**!! The bib and PwCid input cells are not editable unless positively labeled. Please do not open multiple interfaces editing simultaneously.
2. There might be cells that contain an entity mention and a non-entity mention string, e.g., "Bert-large", "Bert with 6 layers frozen". For these cells, you should focous on the entity mention. Thus those two mentions should both be labeled as Method, and linked to https:paperswithcode.com/method/bert.
3. You can scoll down to the end of this page to see indexed bib items.

Annotation Table 1

Verify pre-populated cell types

	Mean ☑ Metric	ceiling ☑ Dataset	floor ☑ Dataset	wall ☑ Dataset	beam ☑ Dataset	column ☑ Dataset	window ☑ Dataset	door ☑ Dataset	table ☑ Dataset	chair ☑ Dataset	sofa ☑ Dataset	bookcase ☑ Dataset	board ☑ Dataset
Armeni et al. [1] ☑ Method	49.93	71.61	88.70	72.86	66.67	91.77	25.92	54.11	46.02	16.15	6.78	54.71	3.91
Seg-Cluster ☑ Method	20.39	43.58	35.52	16.64	12.59	15.90	23.86	15.75	22.63	10.33	3.92	43.33	10.71
SGPN ☑ Method	54.35	79.44	66.29	88.77	77.98	60.71	66.62	56.75	46.90	40.77	6.38	47.61	11.05

Table 1: Results on instance segmentation in S3DIS scenes. The metric is AP(%) with IoU threshold 0.5. To the best of our knowledge, there are no existing instance segmentation methods on point clouds for arbitrary object categories. The result of Armeni et al. [1] is on 3D object detection and...

Submit this sub-table

Annotate attributed sources

Link the Method and Dataset cells to a bib entry (an integer starting from 1), or `-1` if a table cell entity is introduced in this paper, or `0` if its introducing paper is not cited (This can happen for very well-known entities like *Linear Regression*), or `-2` if a cell contains more then one entity.

For Datasets, sometimes you will see 'sub-datasets'. A sub-dataset is a part of a parent-dataset. E.g., ImageNet (the parent), and ImageNet-test (the sub-dataset). **Try to identify the bib entry for the sub-dataset first. If you couldn't find a match, fill in the bib entry for the parent dataset.**

You might need to go back to the paper source page to look for text about the table to figure these out. You can scoll down to the end of this page to see indexed bib items.

In some cases an entity might have more than one introducing papers. In that case, separate the bib entries by `,`. E.g., Smt[1,3] should be labeld as `1,3`.

	Mean	ceiling	floor	wall	beam	column	window	door	table	chair	sofa	bookcase	board
Armeni et	49.93	71.61	88.70	72.86	66.67	91.77	25.92	54.11	46.02	16.15	6.78	54.71	3.91
Seg-Clust	20.39	43.58	35.52	16.64	12.59	15.90	23.86	15.75	22.63	10.33	3.92	43.33	10.71
SGPN	54.35	79.44	66.29	88.77	77.98	60.71	66.62	56.75	46.90	40.77	6.38	47.61	11.05

Table 1: Results on instance segmentation in S3DIS scenes. The metric is AP(%) with IoU threshold 0.5. To the best of our knowledge, there are no existing instance segmentation methods on point clouds for arbitrary object categories. The result of Armeni et al. [1] is on 3D object detection and...

Submit this sub-table

Link cells to entities

Link cells to an entity, if possible, otherwise enter `0`. If a cell is confusing to you, or you think it should be filtered out for the EL task, enter `-2`.

There is a link in each cell to help you. `link to intro paper` links you to the PwC *paper* search page with the introducing paper's title from the previous step. There is a number in each cell indicating the bibliography entry you have identified for the cell. Cliking it will lead you to the bibliography item.

The step for each positively labeled cell should be (note that if you find a match in any step, you are done with this cell and no longer need to go the subsequent steps):

- If it's labeled as Dataset or Dataset&Metric, type in the cell. In cases where sub-datasets appear, please try to search for the most specific match first. **If you do not find a specific match for a sub-dataset, please also try to search for the 'parent dataset' for it. And then make a note that this is a sub-dataset by entering "(sub)PwCName" in the cell.**
- If it's labeled as Method, for each cell you have two things to enter.
  - If it **has an identified attributed paper**, go to the `PwC page of that paper` (using the link) (In rare cases, this might be broken. Please just do manual keyword search with paper title on PwC for those). Then scroll down to the bottom of that PwC page and look for a match in the "Methods" section, click the match and fill in its URL to the first input box if you find a match. Otherwise fills in `0` in the first input box. Then go do the keyword search in the second input box. In the keyword search, please try in the order from **the most-specific key to the least-specific** key. Please also try the long-form if the acronym didn't bring you anything.
  - If it **does not have an identified attributed paper** (i.e., filled in with 0 for the bib task), directly do the keyword serach in the second input box.

	Mean	ceiling	floor	wall	beam	column	window	door	table	chair	sofa	bookcase	board

Armeni et Armeni et	49.93	71.61	88.70	72.86	66.67	91.77	25.92	54.11	46.02	16.15	6.78	54.71	3.91
Seg-Clust Seg-Clust	20.39	43.58	35.52	16.64	12.59	15.90	23.86	15.75	22.63	10.33	3.92	43.33	10.71
SGPN SGPN	54.35	79.44	66.29	88.77	77.98	60.71	66.62	56.75	46.90	40.77	6.38	47.61	11.05

Table 1: Results on instance segmentation in S3DIS scenes. The metric is AP(%) with IoU threshold 0.5. To the best of our knowledge, there are no existing instance segmentation methods on point clouds for arbitrary object categories. The result of Armeni et al. [1] is on 3D object detection and...

Submit this sub-table

Annotation Table 2

Verify pre-populated cell types

	Mean IoU ☑ Metric	Accuracy ☑ Metric
PointNet [31] ☑ Method	49.76	79.66
SGPN ☑ Method	50.37	80.78

Table 4: Results on semantic segmentation in S3DIS scenes. SGPN uses PointNet as baseline. Metric is mean IoU(%) over 13 classes (including clutter)....

Submit this sub-table

Annotate attributed sources

Link the Method and Dataset cells to a bib entry (an integer starting from 1), or `-1` if a table cell entity is introduced in this paper, or `0` if its introducing paper is not cited (This can happen for very well-known entities like *Linear Regression*), or `-2` if a cell contains more then one entity.

For Datasets, sometimes you will see 'sub-datasets'. A sub-dataset is a part of a parent-dataset. E.g., ImageNet (the parent), and ImageNet-test (the sub-dataset). **Try to identify the bib entry for the sub-dataset first. If you couldn't find a match, fill in the bib entry for the parent dataset.**

You might need to go back to the paper source page to look for text about the table to figure these out. You can scoll down to the end of this page to see indexed bib items.

In some cases an entity might have more than one introducing papers. In that case, separate the bib entries by `,`. E.g., Smt[1,3] should be labeld as `1,3`.

	Mean IoU	Accuracy
PointNet [31]	49.76	79.66
SGPN	50.37	80.78

Table 4: Results on semantic segmentation in S3DIS scenes. SGPN uses PointNet as baseline. Metric is mean IoU(%) over 13 classes (including clutter)....

Submit this sub-table

Link cells to entities

Link cells to an entity, if possible, otherwise enter `0`. If a cell is confusing to you, or you think it should be filtered out for the EL task, enter `-2`.

There is a link in each cell to help you. `link to intro paper` links you to the PwC *paper* search page with the introducing paper's title from the previous step. There is a number in each cell indicating the bibliography entry you have identified for the cell. Cliking it will lead you to the bibliography item.

The step for each positively labeled cell should be (note that if you find a match in any step, you are done with this cell and no longer need to go the subsequent steps):

- If it's labeled as Dataset or Dataset&Metric, type in the cell. In cases where sub-datasets appear, please try to search for the most specific match first. **If you do not find a specific match for a sub-dataset, please also try to search for the 'parent dataset' for it. And then make a note that this is a sub-dataset by entering "(sub)PwCName" in the cell.**
- If it's labeled as Method, for each cell you have two things to enter.
  - If it **has an identified attributed paper**, go to the `PwC page of that paper` (using the link) (In rare cases, this might be broken. Please just do manual keyword search with paper title on PwC for those). Then scroll down to the bottom of that PwC page and look for a match in the "Methods" section, click the match and fill in its URL to the first input box if you find a match. Otherwise fills in `0` in the first input box. Then go do the keyword search in the second input box. In the keyword search, please try in the order from **the most-specific key to the least-specific** key. Please also try the long-form if the acronym didn't bring you anything.
  - If it **does not have an identified attributed paper** (i.e., filled in with 0 for the bib task), directly do the keyword serach in the second input box.

	Mean IoU	Accuracy
PointNet [31] PointNet [31]	49.76	79.66
SGPN SGPN	50.37	80.78

Table 4: Results on semantic segmentation in S3DIS scenes. SGPN uses PointNet as baseline. Metric is mean IoU(%) over 13 classes (including clutter)....

Submit this sub-table

Annotation Table 3

Verify pre-populated cell types

	Mean Metric	air- plane Dataset	bag Dataset	cap Dataset	car Dataset	chair Dataset	head phone Dataset	guitar Dataset	knife Dataset	lamp Dataset	laptop Dataset	motor Dataset	mug Dataset	pistol Dataset	rocket Dataset	skate board Dataset	table Dataset
[33] Method	84.6	80.4	80.9	60.0	76.8	88.1	83.7	90.2	82.6	76.9	94.7	68.0	91.2	82.1	59.9	78.2	87.5
SGPN Method	85.8	80.4	78.6	78.8	71.5	88.6	78.0	90.9	83.0	78.8	95.8	77.8	93.8	87.4	60.1	92.3	89.4

Table 7: Semantic segmentation results on ShapeNet part dataset. Metric is mean IoU(%) on points....

Submit this sub-table

Annotate attributed sources

Link the Method and Dataset cells to a bib entry (an integer starting from 1), or `-1` if a table cell entity is introduced in this paper, or `0` if its introducing paper is not cited (This can happen for very well-known entities like *Linear Regression*), or `-2` if a cell contains more then one entity.

For Datasets, sometimes you will see 'sub-datasets'. A sub-dataset is a part of a parent-dataset. E.g., ImageNet (the parent), and ImageNet-test (the sub-dataset). **Try to identify the bib entry for the sub-dataset first. If you couldn't find a match, fill in the bib entry for the parent dataset.**

You might need to go back to the paper source page to look for text about the table to figure these out. You can scoll down to the end of this page to see indexed bib items.

In some cases an entity might have more than one introducing papers. In that case, separate the bib entries by `,`. E.g., Smt[1,3] should be labeld as `1,3`.

	Mean	air- pla	bag	cap	car	chair	head p	guitar	knife	lamp	laptop	motor	mug	pistol	rocket	skate b	table
[33]	84.6	80.4	80.9	60.0	76.8	88.1	83.7	90.2	82.6	76.9	94.7	68.0	91.2	82.1	59.9	78.2	87.5
SGPN	85.8	80.4	78.6	78.8	71.5	88.6	78.0	90.9	83.0	78.8	95.8	77.8	93.8	87.4	60.1	92.3	89.4

Table 7: Semantic segmentation results on ShapeNet part dataset. Metric is mean IoU(%) on points....

Submit this sub-table

Link cells to entities

Link cells to an entity, if possible, otherwise enter `0`. If a cell is confusing to you, or you think it should be filtered out for the EL task, enter `-2`.

There is a link in each cell to help you. `link to intro paper` links you to the PwC *paper* search page with the introducing paper's title from the previous step. There is a number in each cell indicating the bibliography entry you have identified for the cell. Cliking it will lead you to the bibliography item.

The step for each positively labeled cell should be (note that if you find a match in any step, you are done with this cell and no longer need to go the subsequent steps):

- If it's labeled as Dataset or Dataset&Metric, type in the cell. In cases where sub-datasets appear, please try to search for the most specific match first. **If you do not find a specific match for a sub-dataset, please also try to search for the 'parent dataset' for it. And then make a note that this is a sub-dataset by entering "(sub)PwCName" in the cell.**
- If it's labeled as Method, for each cell you have two things to enter.
  - If it **has an identified attributed paper**, go to the `PwC page of that paper` (using the link) (In rare cases, this might be broken. Please just do manual keyword search with paper title on PwC for those). Then scroll down to the bottom of that PwC page and look for a match in the "Methods" section, click the match and fill in its URL to the first input box if you find a match. Otherwise fills in `0` in the first input box. Then go do the keyword search in the second input box. In the keyword search, please try in the order from **the most-specific key to the least-specific** key. Please also try the long-form if the acronym didn't bring you anything.
  - If it **does not have an identified attributed paper** (i.e., filled in with 0 for the bib task), directly do the keyword serach in the second input box.

	Mean	air- pla	bag	cap	car	chair	head p	guitar	knife	lamp	laptop	motor	mug	pistol	rocket	skate b	table
[33] [33]	84.6	80.4	80.9	60.0	76.8	88.1	83.7	90.2	82.6	76.9	94.7	68.0	91.2	82.1	59.9	78.2	87.5
SGPN SGPN	85.8	80.4	78.6	78.8	71.5	88.6	78.0	90.9	83.0	78.8	95.8	77.8	93.8	87.4	60.1	92.3	89.4

Table 7: Semantic segmentation results on ShapeNet part dataset. Metric is mean IoU(%) on points....

Submit this sub-table

1. I. Armeni, O. Sener, A. R. Zamir, H. Jiang, I. Brilakis, M. Fischer, and S. Savarese. 3d semantic parsing of large-scale indoor spaces. In CVPR, 2016.
2. V. Badrinarayanan, A. Kendall, and R. Cipolla. Segnet: A deep convolutional encoder-decoder architecture for image segmentation. TPAMI, 2017.
3. L. Bertinetto, J. Valmadre, J. F. Henriques, A. Vedaldi, and P. H. S. Torr. Fully-convolutional siamese networks for object tracking. In ECCV Workshops, 2016.
4. A. X. Chang, T. Funkhouser, L. Guibas, P. Hanrahan, Q. Huang, Z. Li, S. Savarese, M. Savva, S. Song, H. Su, et al. Shapenet: An information-rich 3d model repository. arXiv preprint arXiv:1512.03012, 2015.
5. X. Chen, H. Ma, J. Wan, B. Li, and T. Xia. Multi-view 3d object detection network for autonomous driving. In CVPR, 2017.
6. S. Chopra, R. Hadsell, and Y. LeCun. Learning a similarity metric discriminatively, with application to face verification. In CVPR, 2005.
7. A. Dai, A. X. Chang, M. Savva, M. Halber, T. Funkhouser, and M. Niessner. Scannet: Richly-annotated 3d reconstructions of indoor scenes. In CVPR, 2017.
8. A. Dai, C. R. Qi, and M. Niessner. Shape completion using 3d-encoder-predictor cnns and shape synthesis. In CVPR, 2017.
9. J. Dai, K. He, Y. Li, S. Ren, and J. Sun. Instance-sensitive fully convolutional networks. In ECCV, 2016.
10. J. Dai, K. He, and J. Sun. Instance-aware semantic segmentation via multi-task network cascades. In CVPR, 2016.

11. Z. Deng and L. J. Latecki. Amodal detection of 3d objects: Inferring 3d bounding boxes from 2d ones in rgb-depth images. In CVPR, 2017.
12. A. Frome, Y. Singer, F. Sha, and J. Malik. Learning globally-consistent local distance functions for shape-based image retrieval and classification. In ICCV, 2007.
13. C.-Y. Fu, W. Liu, A. Ranga, A. Tyagi, and A. C. Berg. Dssd: Deconvolutional single shot detector. arXiv preprint arXiv:1701.06659, 2017.
14. R. Girshick. Fast r-cnn. In ICCV, 2015.
15. R. Girshick, J. Donahue, T. Darrell, and J. Malik. Rich feature hierarchies for accurate object detection and semantic segmentation. In CVPR, 2014.
16. X. Han, T. Leung, Y. Jia, R. Sukthankar, and A. C. Berg. Matchnet: Unifying feature and metric learning for patch-based matching. In CVPR, 2015.
17. X. Han, Z. Li, H. Huang, E. Kalogerakis, and Y. Yu. High-resolution shape completion using deep neural networks for global structure and local geometry inference. In ICCV, 2017.
18. K. He, G. Gkioxari, P. Dollar, and R. Girshick. Mask r-cnn. In ICCV, 2017.
19. D. Kingma and J. Ba. Adam: A method for stochastic optimization. arXiv preprint arXiv:1412.6980, 2014.
20. G. Koch, R. Zemel, and R. Salakhutdinov. Siamese neural networks for one-shot image recognition. In ICML Deep Learning Workshop, 2015.
21. A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. In NIPS, 2012.
22. L. Leal-Taixe, C. Canton-Ferrer, and K. Schindler. Learning by tracking: siamese cnn for robust target association. CVPR DeepVision Workshops, 2016.
23. Y. Li, H. Qi, J. Dai, X. Ji, and Y. Wei. Fully convolutional instance-aware semantic segmentation. In CVPR, 2017.
24. T.-Y. Lin, P. Dollar, R. Girshick, K. He, B. Hariharan, and S. Belongie. Feature pyramid networks for object detection. In CVPR, 2017.
25. T.-Y. Lin, P. Goyal, R. Girshick, K. He, and P. Dollar. Focal loss for dense object detection. In CVPR, 2017.
26. W. Liu, D. Anguelov, D. Erhan, C. Szegedy, S. Reed, C.-Y. Fu, and A. C. Berg. SSD: Single shot multibox detector. In ECCV, 2016.
27. D. Maturana and S. Scherer. Voxnet: A 3d convolutional neural network for real-time object recognition. In IROS, 2015.
28. A. Newell and J. Deng. Associative embedding: End-to-end learning for joint detection and grouping. In NIPS, 2016.
29. P. O. Pinheiro, R. Collobert, and P. Dollar. Learning to segment object candidates. In NIPS, 2015.
30. P. O. Pinheiro, T.-Y. Lin, R. Collobert, and P. Dollar. Learning to refine object segments. In ECCV, 2016.
31. C. R. Qi, H. Su, K. Mo, and L. J. Guibas. Pointnet: Deep learning on point sets for 3d classification and segmentation. CVPR, 2017.
32. C. R. Qi, H. Su, M. Niessner, A. Dai, M. Yan, and L. J. Guibas. Volumetric and multi-view cnns for object classification on 3d data. In CVPR, 2016.
33. C. R. Qi, L. Yi, H. Su, and L. J. Guibas. Pointnet++: Deep hierarchical feature learning on point sets in a metric space. In NIPS, 2017.
34. X. Qi, R. Liao, J. Jia, S. Fidler, and R. Urtasun. 3d graph neural networks for rgbd semantic segmentation. In CVPR, 2017.
35. J. Redmon, S. Divvala, R. Girshick, and A. Farhadi. You only look once: Unified, real-time object detection. In CVPR, 2016.
36. J. Redmon and A. Farhadi. Yolo9000: Better, faster, stronger. In CVPR, 2017.
37. S. Ren, K. He, R. Girshick, and J. Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. In NIPS, 2015.
38. Z. Ren and E. B. Sudderth. Three-dimensional object detection and layout prediction using clouds of oriented gradients. In CVPR, 2016.
39. G. Riegler, A. O. Ulusoy, and A. Geiger. Octnet: Learning deep 3d representations at high resolutions. In CVPR, 2017.
40. N. Silberman, D. Hoiem, P. Kohli, and R. Fergus. Indoor segmentation and support inference from rgbd images. ECCV, 2012.
41. E. Simo-Serra, E. Trulls, L. Ferraz, I. Kokkinos, P. Fua, and F. Moreno-Noguer. Discriminative learning of deep convolutional feature point descriptors. In ICCV, 2015.
42. S. Song and J. Xiao. Sliding shapes for 3d object detection in depth images. In ECCV, 2014.
43. S. Song and J. Xiao. Deep Sliding Shapes for amodal 3D object detection in RGB-D images. In CVPR, 2016.
44. S. Song, F. Yu, A. Zeng, A. X. Chang, M. Savva, and T. Funkhouser. Semantic scene completion from a single depth image. arXiv preprint arXiv:1611.08974, 2016.
45. M. Tatarchenko, A. Dosovitskiy, and T. Brox. Octree generating networks: Efficient convolutional architectures for high-resolution 3d outputs. In ICCV, 2017.
46. P.-S. Wang, Y. Liu, Y.-X. Guo, C.-Y. Sun, and X. Tong. O-cnn: Octree-based convolutional neural networks for 3d shape analysis. ACM Transactions on Graphics (TOG), 2017.
47. W. Wang, Q. Huang, S. You, C. Yang, and U. Neumann. Shape inpainting using 3d generative adversarial network and recurrent convolutional networks. In ICCV, 2017.

48. K. Q. Weinberger and L. K. Saul. Distance metric learning for large margin nearest neighbor classification. Journal of Machine Learning Research, 2009.

49. Z. Wu, S. Song, A. Khosla, F. Yu, L. Zhang, X. Tang, and J. Xiao. 3d shapenets: A deep representation for volumetric shapes. In CVPR, 2015.

50. D. Yi, Z. Lei, S. Liao, and S. Z. Li. Deep metric learning for person re-identification. In ICPR, 2014.

51. L. Yi, V. G. Kim, D. Ceylan, I. Shen, M. Yan, H. Su, A. Lu, Q. Huang, A. Sheffer, L. Guibas, et al. A scalable active framework for region annotation in 3d shape collections. ACM Transactions on Graphics (TOG), 2016.